

# CS 555: DISTRIBUTED SYSTEMS

## [REPLICATION & CONSISTENCY]

Shrideep Pallickara  
Computer Science  
Colorado State University

October 29, 2019

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.1

## Frequently asked questions from the previous class survey

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.2

## Topics covered in this lecture

- Replication
- Consistency Models
- Data centric consistency model
  - ▣ Continuous consistency models
  - ▣ Sequential consistency

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.3

## REPLICATION & CONSISTENCY

October 29, 2019

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.4

## What we will look at in our discussions

- Replication
- Consistency
  - ▣ Models
  - ▣ Client models
  - ▣ Protocols
- Eventual Consistency
- Brewer's CAP Theorem

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.5

## Why are these inter-related topics important?

- Performance
  - Correctness
  - Failure to account for interactions between these issues?
    - ▣ Poor performance
    - ▣ Inaccurate results
- } The holy grail of demonstrable incompetency in systems development!

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.6

## REPLICATION

October 29, 2019

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.7

### Rationale for replication

- Reliability
- Availability
- Performance

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.8

## Rationale for replication: Reliability

- Replication as a safeguard against **failures**
- Protection against data **corruptions**
- File System example:
  - ▣ **3** copies
  - ▣ If one fails, process can choose from the other two
  - ▣ Read/write performed on each copy
    - At least 2 of the reads must **concur**
    - Protects against a failing write

} Data  
corruptions

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.9

## Rationale for replication: Increased Availability

- Users require services to be **highly available**
  - ▣ Proportion of time when service is **accessible with reasonable response times** should be close to 100%
- Factors relevant to high-availability
  - ▣ Delays due to pessimistic concurrency control
  - ▣ Server failures
  - ▣ Network partitions and disconnected operations

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.10

## Replication maintains availability despite server failures [1/2]

- Data is replicated at failure independent servers
- Client software should be able to access data at an alternative server if default server fails

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.11

## Replication maintains availability despite server failures [2/2]

- If each of the  $n$  servers has an independent probability  $p$  of failing or becoming unreachable
- The availability of an object stored at each of these servers?
  - 1- *probability(all servers fail or are unreachable)*
  - $1 - p^n$

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.12

## Replication maintains availability despite server failures: Example

- There is a 5% probability of independent server failures?
- There are two servers
  - ▣ Availability is  $1 - p^n$
  - ▣  $1 - (0.05)^2 = 1 - 0.0025 = 99.75\%$

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.13

## Rationale for replication: Performance

- Ability to scale with **numbers**
  - ▣ Processes access data managed by a server
  - ▣ Replicate server; distribute work
- Ability to scale with **geographical area**
  - ▣ Place copy of data in *proximity* of processes using it
  - ▣ Time to access service decreases
    - **Perceived performance** improves

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.14

## But replication exacts a price ...

- A client may perceive better performance but ...
  - ▣ More *network bandwidth* needed
    - To keep replicas in sync
- **Consistency** problems
  - ▣ When a copy is modified, it becomes *different*
  - ▣ Modifications have to be made on **all** copies

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.15

## Replication Costs: *When and how* modifications must be made to copies

- Fetching a page from a remote Web server
  - ▣ OBJECTIVE: Improving access times
- Web browsers locally **cache** a web page
  - ▣ If user requests the same page
    - Returned from cache
    - User is happy with the load times
  - ▣ What if user always wants the latest copy?

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.16



## Simple solutions to the stale copy problem

- ① **Don't cache** web page
  - ▣ If there is no nearby replica, performance is poor
  - ▣ Also, what if the page does not change that often?
- ② Let server **invalidate/update** caches
  - ▣ Server must track all caches
  - ▣ Degrades server performance

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.17

## Replication as a scaling technique

- ▣ Placing data copies close to processes
  - ▣ Improves access times
  - ▣ Distributes work
- ▣ Potential problems ...

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.18

## Replication for scaling: Network bandwidth

- Process **P** accesses a replica **N** times per second
- Replica is itself updated **M** times per second
- If **N** << **M** ?
  - ▣ Several updated versions of replica *never* accessed
  - ▣ *Network traffic* to install those versions: wasted!
  - ▣ Perhaps installing a replica was not a good idea?

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.19

## Replication for scaling: Consistency issues

- Consistency might itself be subject to **scaling problems**
- Collection of copies is consistent when *all* copies are the same
  - ▣ Read on any copy returns the *same* result
  - ▣ Updates propagated to all copies before the next operation?
    - **Tight consistency**

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.20

## Consistency issues in replication

- Update performed at all copies as an **atomic operation**
  - ▣ Transaction
- Implementing atomicity with large number of replicas is difficult
  - ▣ May be dispersed on a WAN
  - ▣ Operations **cannot** complete quickly

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.21

## Other things that replicas need to agree on ...

- Replicas must agree on **when** operation must be performed locally
- Replicas need to decide on **ordering**
  - ▣ Lamport timestamps
  - ▣ Coordinator assigned order

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.22

## The Replication Dilemma

- Alleviating scalability issues
  - ▣ Replication and caching: Improves performance
- Keeping copies consistent?
  - ▣ Requires *global* synchronization
  - ▣ *Costly* in terms of performance
    - Time
    - Network bandwidth

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.23

## DATA CENTRIC CONSISTENCY MODELS

October 29, 2019

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.24

## Data centric consistency models


- Consistency is in the context of read/write operation on **distributed, shared** data
  - ▣ Memory
  - ▣ Database
  - ▣ File systems
- The broader term **data store** is more commonly used

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.25

## Distributed data store



October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.26

## Consistency model

- **Contract** between processes and the data store
- If processes agree to obey certain **rules**
  - ▣ Data store works correctly

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.27

## Consistency that we intuitively expect

- Process performing a **read** on a data item
  - ▣ Expects value to show results of **last write** operation on that item
- Without a global clock?
  - ▣ Difficult to define **which** write was the last one

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.28

## We thus need to provide other definitions ... consistency models

- Each model **restricts values** that a **read** operation on a data item can return
- Models with the greatest restrictions
  - ▣ Easiest to use
- Models with minor restrictions
  - ▣ Difficult to use
- Easy-to-use models **do not** perform as well as difficult ones

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.29

## Loosening of consistency

- Needed for efficiency and performance
- No **general rules** however
  - ▣ **Tolerance depends on the application**

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.30

## CONTINUOUS CONSISTENCY MODELS

October 29, 2019

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.31

### Continuous consistency

- **Three axes** for defining inconsistencies
- **Deviations** between replicas in terms of
  - Numerical values
  - Staleness between replicas
  - Ordering of update operations
- Deviations form **continuous consistency** ranges

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.32



## Example of using continuous consistency models: Stock prices

- Two copies of a stock should not deviate by more than 2 cents.
  - ▣ **Absolute numerical deviation**
- Two copies do not deviate by more than 0.5%
  - ▣ **Relative numerical deviation**
- If stock goes up and one replica is updated
  - ▣ If change **does not** violate specified deviations?
    - Replicas are considered consistent

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.33

## Numerical and Staleness deviations

- Numerical deviation can also be expressed in terms of number of **updates**
  - ▣ Applied at a replica, but not seen by other replicas
- **Staleness** deviations
  - ▣ Last time a replica was updated
  - ▣ Replica can provide **old data** as long as it is not too old
    - Weather reports

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.34

## Ordering of updates may also be allowed to be different

- Within a certain **bound**
- Updates applied *tentatively* at local copy
  - ▣ Need global agreement with all replicas
- Before an update becomes *permanent* it
  - ▣ Might be rolled back
  - ▣ Applied in a different order

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.35

## CONSISTENCY UNIT (CONIT)

October 29, 2019

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.36

## Consistency Unit: conit

- Specifies *unit* over which consistency is to be **measured**
- Examples
  - Record representing a stock
  - Weather report

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.37

## Looking at the conit a little closer: Example with 2 replicas

- Each replica maintains a 2D vector clock
- Operation carried out by *replica i* at (its) **logical time  $t : \langle t, i \rangle$**
- Example conit contains data items  $x$  and  $y$

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.38

## Tracking consistency deviations: Conit items $x$ and $y$ are initialized to 0

Replica A

Conit: $x=6, y=3$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 8, A \rangle$	$y = y+2$	$[y = 2]$
$\langle 12, A \rangle$	$y = y+1$	$[y = 3]$
$\langle 14, A \rangle$	$x = y \times 2$	$[x = 6]$

Committed Operation: 

Replica B

Conit: $x=2, y=5$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 10, B \rangle$	$y = y+5$	$[y = 5]$

Tentative Operation: 

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.39

## Vector Clocks at each replica

Replica A

Conit: $x=6, y=3$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 8, A \rangle$	$y = y+2$	$[y = 2]$
$\langle 12, A \rangle$	$y = y+1$	$[y = 3]$
$\langle 14, A \rangle$	$x = y \times 2$	$[x = 6]$

Vector clock A = (15, 5)

Replica B

Conit: $x=2, y=5$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 10, B \rangle$	$y = y+5$	$[y = 5]$

Vector clock B = (0, 11)

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.40

## Order deviations are the number of tentative operations at each replica

Replica A

Conit: $x=6, y=3$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 8, A \rangle$	$y = y+2$	$[y = 2]$
$\langle 12, A \rangle$	$y = y+1$	$[y = 3]$
$\langle 14, A \rangle$	$x = y \times 2$	$[x = 6]$

Order Deviation = 3

Replica B

Conit: $x=2, y=5$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 10, B \rangle$	$y = y+5$	$[y = 5]$

Order Deviation = 2

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.41

## Numerical deviations in our example

- Numerical deviation here is the number of **unseen updates** from the other replica
- Weight of this deviation at replica **A** is the maximum **difference** between
  - ▣ Committed values of conit at **A**
  - ▣ Result from operations at **B** not seen by **A**

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.42

## Quantifying the numerical deviations at each replica

Replica A

Conit: $x=6, y=3$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 8, A \rangle$	$y = y+2$	$[y = 2]$
$\langle 12, A \rangle$	$y = y+1$	$[y = 3]$
$\langle 14, A \rangle$	$x = y \times 2$	$[x = 6]$

Unseen Updates = 1  
 Weight =  $\text{Max}[\text{diff}(2,2), \text{diff}(0,5)]$   
 = 5

Replica B

Conit: $x=2, y=5$		
Operation	Result	
$\langle 5, B \rangle$	$x = x+2$	$[x = 2]$
$\langle 10, B \rangle$	$y = y+5$	$[y = 5]$
Note: B's committed value is (0,0)		

Unseen updates = 3  
 Weight =  $\text{Max}[\text{diff}(0,6), \text{diff}(0,3)]$   
 = 6

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.43

## Tradeoffs between fine grained and coarse grained conits

- If conit represents a **lot** of data
  - ▣ Updates aggregated for all data in conit
  - ▣ Replicas become **inconsistent sooner**
- If conit is smaller
  - ▣ Fewer updates needed
  - ▣ **Total number** of conits to be managed goes up

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.44

## Before we put conits to practical use two things need to happen

- **Protocols** to enforce consistency
- Developers **specify** consistency requirements
  - Difficult!

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.45

## Conits are declared alongside updates

AffectsConit(**ConitQ**, 1, 1)  
append message **m** to queue **Q**

- Appending message **M** to queue **Q** belongs to a conit named **ConitQ**

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.46

## Conits are declared alongside reads

`DependsOnConit(ConitQ, 4, 0, 60)`  
read message `m` from the head of queue `Q`

- Numerical deviation: 4
  - ▣ At most 4 unseen updates at other replicas
- Ordering deviation: 0
  - ▣ No tentative local updates
- Staleness deviation: 60 seconds
  - ▣ Check `Q` for staleness periodically

## CONSISTENT ORDERING OF OPERATIONS



## Consistent ordering of operations

- Class of models from **concurrent programming**
- We will look at
  - ▣ Sequential consistency
  - ▣ Causal consistency

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.49

## Sequential consistency: Notations


- Operations of processes depicted along time axis
- Write by a process  $P_i$  to data item  $x$  with value  $a$ 
  - $W_i(x)a$
- Read by a process  $P_i$  of data item  $x$  that returns the value  $b$ 
  - $R_i(x)b$
- All items are initially **NIL**

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.50

## Two processes operating on the same data item



October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.51

## Sequential consistency


- Defined by Lamport
  - ▣ Context: Shared memory in multiprocessor setting
- When processes run concurrently
  - ▣ Any valid interleaving of read/write is acceptable
  - ▣ But all processes **must see the same interleaving**

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.52

## Sequential consistency example




Write operation of P2 appears to be **before** P1  
 This is **acceptable**

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.53

## Sequential consistency: Example



P3 concludes final value is **a**  
 P4 concludes final value is **b**  
**Unacceptable**

October 29, 2019  
 Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
 Dept. Of Computer Science, Colorado State University

L19.54

## Sequential Consistency: Another example

### Process 1

x = 1  
print(y,z)

### Process 2

y = 1  
Print(x,z)

### Process 3

z = 1  
Print(x,y)

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.55

## Multiple interleaved sequences are possible

- With 6 statements there are
  - ▣ 6! possibilities = 720
  - ▣ Some of these **violate program order**
- 120 (5!) sequences begin with x=1
  - ▣ Half print(x,z) before y=1
    - ▣ Half print(x,y) before z=1
      - Only 1/4 or 30 are valid
- Similarly, there are 30 that start with y=1, z=1
  - ▣ Total of 90 valid execution sequences

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.56

## Different, but valid interleaving of the statements

**Signature** is the concatenation of the outputs of P1, P2 and P3

```
x = 1
print(y,z)
y = 1
print(x,z)
z = 1
print(x,y)
```

**Prints:** 001011  
**Signature:** 001011

```
x = 1
y = 1
print(x,z)
print(y,z)
z = 1
print(x,y)
```

**Prints:** 101011  
**Signature:** 101011

```
y = 1
z = 1
print(x,y)
print(x,z)
x = 1
print(y,z)
```

**Prints:** 010111  
**Signature:** 110101

```
y = 1
x = 1
z = 1
print(x,z)
print(y,z)
print(x,y)
```

**Prints:** 111111  
**Signature:** 111111

## Contract between processes and shared data store

- Processes must accept **all valid results**
- Must work if any of them occurs

## Invalid sequences in signature patterns

- 000000?
  - ▣ Print statements ran before assignments
  - ▣ **Violates** program order
- 001001?
  - ▣ {00} y and z were 0 when **P1** did its printing
    - **P1** executes its statements *before* **P2** and **P3** start
  - ▣ {10} **P2** ran after **P1** started, but before **P3** started
  - ▣ {01} **P3** must complete *before* **P1** starts
    - Not possible!

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.59

## The contents of this slide-set are based on the following references

- *Distributed Systems: Concepts and Design*. George Coulouris, Jean Dollimore, Tim Kindberg, Gordon Blair. 5th Edition. Addison Wesley. ISBN: 978-0132143011. [Chapter 6, 7]
- *Distributed Systems: Principles and Paradigms*. Andrew S. Tanenbaum and Maarten Van Steen. 2nd Edition. Prentice Hall. ISBN: 0132392275/978-0132392273. [Chapter 4, 18]

October 29, 2019  
Professor: SHRIDEEP PALICKARA

CS555: Distributed Systems [Fall 2019]  
Dept. Of Computer Science, Colorado State University

L19.60